

# Actor Free critic Update for Off-policy and Offline learning

## ABSTRACT

This work investigates Actor-Free critic Updates (AFU), a novel reinforcement learning algorithm that decouples critic updates from the actor through value-advantage decomposition, comparing its performance against traditional actor-critic methods in both off-policy learning with randomly sampled data and offline-to-online transition scenarios. Our experiments across multiple continuous control environments demonstrate that while AFU does not consistently outperform traditional methods in pure off-policy learning with random data, it exhibits superior stability during the critical transition from offline to online learning, suggesting significant potential for applications where sample efficiency and smooth deployment from pretraining are essential.

## INTRODUCTION

Reinforcement learning algorithms for continuous control have seen remarkable advances with actor-critic architectures becoming the standard approach. These methods combine policy gradient techniques with value-based learning to enable sample-efficient optimization in continuous action spaces. However, traditional actor-critic methods like Deep Deterministic Policy Gradient and Soft Actor-Critic face important limitations in their ability to learn from arbitrary data distributions and other offline methods like Implicit Q-Learning and Calibrated Q-Learning show limitations to transition smoothly from offline to online learning.

In this work, we investigate Actor-Free critic Updates, a recently proposed algorithm that fundamentally redesigns the actor-critic architecture. Unlike traditional approaches where critic updates depend on actions sampled by the actor, AFU maintains critic updates completely independent from the actor through value and advantage decomposition with conditional gradient scaling. This structural departure theoretically enables more effective learning from arbitrary data distributions and potentially smoother

transitions between offline and online learning phases.

We conducted comprehensive experiments to test two main hypotheses: first, that AFU can learn more effectively from randomly sampled state-action pairs than traditional algorithms due to its actor-independent critic updates; second, that AFU exhibits greater stability during the transition from offline to online learning due to its approach to solving the max-Q problem. Our experiments spanned multiple continuous control environments of varying complexity, from simple tasks like CartPole to more complex scenarios like Lunar Lander, providing a thorough evaluation of these hypotheses.

## RELATED WORK

Reinforcement learning algorithms for continuous control domains have evolved significantly in recent years. The development of actor-critic architectures has proven particularly effective, with algorithms like DDPG [4] and SAC [5] establishing themselves as standard approaches. These methods combine the advantages of policy gradient methods with value-based learning, enabling sample-efficient learning in continuous action spaces.

DDPG employs a deterministic actor that maximizes a learned Q-function, operating entirely off-policy to improve sample efficiency. However, DDPG presents stability issues, demonstrating sensitivity to hyperparameter choices and exploration strategies. SAC extends this framework by incorporating entropy maximization, encouraging exploration while learning a stochastic policy. The entropy term provides additional stability, allowing SAC to achieve both better sample efficiency and final performance compared to DDPG in many domains.

The challenge of addressing the distribution shift when transitioning from offline to online reinforcement learning has received increasing attention. Offline reinforcement learning methods train agents using previously collected datasets without environment interaction, eliminating exploration costs but introducing optimization difficulties. IQL [1] approaches this problem by completely avoiding direct querying of the learned Q-function with unseen actions during training, using expectile regression to estimate the maximum value of Q-functions. Similarly, CALQL [2] applies a constraint to the conservative Q-learning framework to reduce overestimation bias during offline learning while enabling efficient online fine-tuning. Both methods aim to mitigate performance degradation observed when agents trained offline begin interacting with environments directly.

The ability to learn truly off-policy—from data generated by arbitrary or random policies—represents another research direction. Traditional actor-critic methods are categorized as off-policy but often struggle when presented with data significantly different from their current policy distribution. This limitation arises because their critic updates depend on actions sampled by the actor, creating an implicit coupling that restricts genuine off-policy learning. Some algorithms have attempted to address this issue through several methods but the fundamental actor-critic interdependence remains.

AFU [3] introduces a structural departure from previous approaches by maintaining critic updates that remain entirely independent from the actor. Unlike other algorithms derived from Q-learning for continuous control, AFU aims to solve the maximization problem inherent in Q-learning through a mechanism based on value and advantage decomposition, employing conditional gradient scaling. This approach potentially enables more effective learning from arbitrary data distributions without requiring explicit constraints during the critic learning phase.

## PRELEMINARIES

We consider a discounted infinite horizon Markov Decision Problem (MDP) defined as a tuple  $\langle S, A, T, R, \gamma \rangle$ , where  $S$  represents the state space,  $A$  denotes a continuous action space,  $T$  is a stochastic transition function,  $R : S \times A \rightarrow \mathbb{R}$  is a reward function, and  $0 \leq \gamma < 1$  is a discount factor. When an agent performs an action  $a \in A$  in state  $s \in S$ , it transitions to a new state  $s'$  according to the transition probability  $T(s'|s, a)$  and receives a reward  $r = R(s, a)$ . We denote transitions as tuples  $(s, a, r, s')$ .

The goal in reinforcement learning is to find a policy  $\pi : S \rightarrow A$  that maximizes the expected sum of discounted rewards. The optimal Q-function  $Q^*$  is defined as:

$$Q^*(s, a) = \mathbb{E} \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_0 = a, \pi^* \right]$$

where the policy used from  $t = 1$  onwards is  $\pi^*$ , which selects actions optimally in every state. The optimal value function  $V^*$  satisfies  $V^*(s) = \max_{a \in A} (Q^*(s, a))$ .

In deep reinforcement learning, we approximate the value functions using neural networks. We denote by  $V_\varphi$  a function approximators for the value function, and by  $Q_\psi$  a function approximator for the Q-function (the critic), where  $\varphi$ , and  $\psi$  are the parameter vectors of neural networks.

On-policy and off-policy learning represent two distinct approaches in reinforcement learning. In on-policy learning, the agent learns from data collected using its current policy. In contrast, off-policy learning allows the agent to learn from data collected using any policy, including random exploration or previously stored experiences. This distinction is crucial as truly off-policy algorithms can potentially learn more efficiently by reusing diverse experiences.

Offline reinforcement learning takes the off-policy concept further by learning entirely from a fixed dataset of previously collected transitions without any environment interaction during training. When transitioning from offline to online learning (where the agent begins to interact with the environment), algorithms often suffer from distribution shift problems as the learned policy encounters states and actions not represented in the offline dataset.

Actor-critic methods combine policy-based and value-based approaches. The actor (policy network) determines which actions to take, while the critic (value network) evaluates these actions. In traditional actor-critic architectures like SAC and DDPG, the critic’s updates depend on actions sampled by the actor, creating an interdependence between the two components. The AFU algorithm represents a departure from this approach by updating the critic independently of the actor.

The temporal difference (TD) learning used in many algorithms aims to minimize the following loss function for the critic:

$$L_Q(\psi) = \mathbb{E}_{(s,a,r,s') \sim B} \left[ \left( Q_\psi(s,a) - r - \gamma V_\psi(s') \right)^2 \right]$$

where  $B$  represents a mini-batch of transitions sampled from an experience replay buffer.

## METHODS

*Environment.* For our experiments, we utilized the *Gymnasium* framework, a widely adopted benchmark for reinforcement learning

research. Gymnasium provides standardized environments, allowing for reproducible evaluation of algorithmic performance. We primarily focused on continuous control tasks including Pendulum, CartPole (continuous version), LunarLander (continuous version) and MountainCar (continuous version).

To facilitate our research on off-policy learning properties, we extended these environments with additional functionality. Each environment was modified to support direct state manipulation through the implementation of a `_set_state()` method. This modification enables control over the system state, allowing us to sample uniformly from the state-action space during our off-policy learning experiments. It should be noted that such direct state manipulation represents a research tool rather than a practical capability in real-world scenarios, where complete state control is rarely possible.

*Algorithms.* We implemented several state-of-the-art deep reinforcement learning algorithms using PyTorch and the BBRL (BlackBoard Reinforcement Learning) framework. BBRL provides a modular architecture where agents interact through a shared workspace, facilitating the implementation of complex algorithms with clear separation of components. Each algorithm implementation follows the same structure, comprising neural network architectures, training procedures, and evaluation methods.

(*DDPG*) Deep Deterministic Policy Gradient combines the deterministic policy gradient algorithm with deep neural networks. It consists of two primary networks. The actor network implements a deterministic policy  $\pi(s)$  that maps states to specific actions. The critic network evaluates the quality of state-action pairs by approximating the Q-function  $Q(s,a)$ . For stable learning, DDPG employs target networks for both actor and critic, which are updated using soft updates:  $\theta^{\text{target}} \leftarrow \tau\theta + (1 - \tau)\theta^{\text{target}}$ , where  $\tau$  is a small value controlling the update rate. DDPG also uses a replay buffer to store and randomly sample transitions, breaking the correla-

tion between consecutive samples and stabilizing learning.

(*SAC*) Soft Actor-Critic extends the actor-critic architecture by incorporating entropy maximization, encouraging exploration through stochastic policies. Our SAC implementation includes a policy network, a Q-network and a V-network. Unlike DDPG’s deterministic policy, SAC’s policy is stochastic, outputting a Gaussian distribution over actions. The network produces both the mean  $\mu(s)$  and log standard deviation  $\log(\sigma(s))$  of this distribution. Actions are sampled using the reparameterization trick and squashed through a tanh function to bound them. SAC employs two Q-networks to mitigate overestimation bias, a common issue in Q-learning. Both networks approximate  $Q(s, a)$  and the minimum of their predictions is used for updates. The value network estimates the state value  $V(s)$ , which represents the expected future return starting from state  $s$  when following the policy. SAC optimizes the expected return plus the entropy of the policy using a temperature parameter that determines the relative importance of entropy versus reward. This parameter is automatically adjusted during training to achieve a target entropy.

(*AFU*) The Actor-Free critic Updates algorithm represents a significant departure from traditional actor-critic methods. While algorithms like DDPG and SAC update their critic using actions generated by the actor, AFU decouples these components, enabling critic updates that are truly independent of the actor. AFU’s architecture consists of two value networks, two advantage networks, a Q-network and a policy network. The advantage networks estimates the the advantage function  $A(s, a)$ , which represents how much better taking action  $a$  in state  $s$  is compared to the average action.

The key innovation in AFU lies in how it computes the value and advantage functions. Instead of relying on the actor to estimate the maximum value of  $Q(s, a)$  over actions, AFU directly decomposes  $Q(s, a)$  into  $V(s) + A(s, a)$ .

This decomposition, combined with a conditional gradient scaling mechanism, allows AFU to solve the maximization problem in Q-learning without depending on the actor. This mechanism allows AFU to maintain stable learning by adaptively adjusting the gradient flow depending on whether the current estimate falls short of the target.

(*IQL*) Implicit Q-Learning approaches offline reinforcement learning by avoiding direct querying of the Q-function with unseen actions during training. This algorithm uses expectile regression to estimate the maximum value of Q-functions without requiring access to the optimal action. IQL operates on three main components: a value network, two Q-networks to reduce overestimation bias, and a policy network. The value function is trained using expectile regression at a specific quantile  $\tau$ , enabling it to approximate the maximum of the Q-function distribution. The policy is then extracted through advantage-weighted regression, allowing it to favor actions with higher advantages while remaining within the data distribution. This approach effectively addresses the distribution shift problem in offline RL by never evaluating actions outside the training dataset, making it well-suited for purely offline learning scenarios but potentially limiting its adaptability during transitions to online learning.

(*Cal-QL*) Calibrated Q-Learning builds upon conservative Q-learning approaches for offline reinforcement learning while enabling smoother transition to online learning. CALQL employs a Q-network, a target network, and a policy network similar to SAC’s stochastic policy. What distinguishes CALQL is its use of reference values precomputed from the dataset using Monte Carlo returns, which serve as a calibration mechanism. The learning objective combines a standard temporal difference error with a conservative regularization term that explicitly penalizes out-of-distribution actions. This calibration helps prevent both underestimation and overestimation of Q-values, addressing a key challenge in offline-to-online transition where traditional

conservative approaches tend to be overly pessimistic. CALQL’s balanced approach allows it to maintain sufficient conservatism during offline learning while facilitating effective exploration when transitioning to online learning.

*(Off-policy learning experimental setup)* Our first experiment aims to evaluate whether AFU can truly learn from randomly generated data, a capability that would distinguish it from traditional off-policy algorithms like DDPG and SAC. The hypothesis stems from the theoretical properties of AFU’s critic update mechanism, which, unlike SARSA-based algorithms, does not rely on the actor’s improvement to approximate the max-Q operation.

Traditional actor-critic algorithms like SAC and DDPG are structurally closer to SARSA than to Q-learning. Despite being categorized as off-policy, they depend on the actor to generate actions for the critic’s updates, creating an implicit coupling. In contrast, AFU’s value and advantage decomposition, combined with conditional gradient scaling, allows it to compute critic updates independently of the actor, potentially enabling true off-policy learning.

To test this hypothesis, we designed a protocol where instead of collecting experience through environment interaction with the current policy, we randomly sample states from the observation space and actions from the action space. This sampling is uniform, representing the extreme case of off-policy data with no correlation to any learning policy. If AFU can converge under these conditions while traditional algorithms struggle, it would validate its theoretical advantage in truly off-policy learning.

*(Offline-to-online transition experimental setup)* Our second experiment investigates the stability of algorithms during the transition from offline to online learning. This transition presents a significant challenge due to the distribution shift between the fixed offline dataset and the data generated by the current policy during online learning.

Offline reinforcement learning algorithms often employ conservatism to prevent overestimation of out-of-distribution actions. While this conservatism is beneficial during offline learning, it can hinder exploration during subsequent online learning. Algorithms like IQL and CALQL address this through different mechanisms: IQL uses expectile regression to estimate maximum values without querying unseen actions, while CALQL explicitly constrains Q-values for out-of-distribution actions.

We hypothesize that AFU may exhibit superior stability during this transition due to its ability to estimate maximum Q-values without relying on the actor. This capability potentially reduces the need for explicit conservatism, allowing AFU to adapt more smoothly when transitioning to online learning.

To test this hypothesis, we generated datasets from policies at different stages of training, capturing trajectories at fixed intervals during policy learning. We then trained algorithms offline on these datasets before transitioning to online learning. The primary evaluation metric is the area between the learning curve and the maximum stable performance level following the transition point, which quantifies how quickly and stably an algorithm recovers during the online phase.

## RESULTS

*Off-policy Learning.* In this section, we present and analyze the results of our off-policy learning experiments across various continuous control environments. These experiments test our hypothesis that AFU’s structural independence between critic and actor updates would provide advantages in learning from purely random data compared to traditional actor-critic algorithms.

We first examine the results on three environments of increasing complexity: CartPole, MountainCar, and Pendulum. In the CartPole environment ( Fig. 1), all three algorithms demonstrated successful convergence when trained on randomly sampled state-action pairs. This outcome

contradicts our initial hypothesis, as we expected traditional actor-critic methods like SAC and DDPG to struggle with entirely random data. The histogram analysis (Fig. 8) confirms that by the end of training, all algorithms converged to similar near-optimal reward distributions.

An interesting observation in the CartPole results is a temporary performance drop around 50,000 training steps for both AFU and SAC. This pattern may be attributed to their entropy-maximizing approach, which encourages exploration of potentially promising but ultimately suboptimal regions of the state-action space. After this exploratory phase, both algorithms recover and converge to robust solutions.




Fig. 1. – Evolution of training performance for SAC, AFU, and DDPG algorithms in the CartPole environment under off-policy learning with uniformly sampled state-action pairs. The curves represent interquartile means of returns over 200,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. Performance evaluations were conducted every 100 training steps by deploying the current policy for 10 complete episodes.

The MountainCar environment presents a more challenging task due to its deceptive reward structure, where energy expenditure necessary

to climb the hill is penalized unless the flag is reached. Fig. 2 shows that all three algorithms successfully converged under random sampling, while none achieved optimal performance in on-policy training (Fig. 14). This difference highlights how uniform sampling across the state-action space can overcome the exploration challenges posed by deceptive gradients. The uniform sampling enables discovery of the rare rewarding state-action sequences that might never be encountered through standard exploration policies.




Fig. 2. – Evolution of training performance for SAC, AFU, and DDPG algorithms in the MountainCar environment under off-policy learning with uniformly sampled state-action pairs. The curves represent interquartile means of returns over 200,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. Performance evaluations were conducted every 100 training steps by deploying the current policy for 10 complete episodes.

The Pendulum environment provides particularly revealing insights due to its initial state dependency. In on-policy training (Fig. 16), SAC and DDPG quickly converged to near-optimal solutions, while AFU displayed significant performance variance. Analysis of this variance revealed that AFU successfully learned strategies for favorable initial conditions but struggled with

cases where the pendulum starts at the bottom position with zero momentum, failing to learn the necessary swinging motion.




Fig. 3. – Evolution of training performance for SAC, AFU, and DDPG algorithms in the Pendulum environment under off-policy learning with uniformly sampled state-action pairs. The curves represent interquartile means of returns over 200,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. Performance evaluations were conducted every 100 training steps by deploying the current policy for 10 complete episodes.

Surprisingly, in off-policy training with random sampling (Fig. 3), AFU converged rapidly alongside SAC and DDPG, showing none of the difficulties observed in on-policy training. The reward distributions (Fig. 10) confirm that all algorithms achieved consistently near-optimal performance.

The value function visualizations for Pendulum (Fig. 22, Fig. 23, and Fig. 24) provide deeper insights into the learned policies. All three algorithms under off-policy training developed similar spiral-structured value functions, characteristic of optimal pendulum control. This contrasts with AFU’s on-policy visualization (Fig. 25), which shows a notably different value estimation

at the critical bottom position, incorrectly assigning high value to states where the pendulum is at the bottom with zero momentum.

Our results require us to reconsider the mechanisms behind off-policy learning in deep reinforcement learning. While our hypothesis focused on AFU’s ability to solve the maximum Q-value problem independently of the actor, our experiments revealed a more nuanced picture of critic updating mechanisms.

In reinforcement learning with continuous action spaces, the critic update requires estimating  $V(s') = \max_a Q(s', a')$  for the target value calculation. This creates the « max Q problem. » Traditional actor-critic methods approximate this by using  $V(s') \approx Q(s', \pi(s'))$ , creating an implicit dependency on the actor’s quality. AFU instead decomposes  $Q(s, a)$  into  $V(s) + A(s, a)$  and uses conditional gradient scaling to solve the maximization problem directly. This technical approach to solving the max Q problem was a key motivation for our study, as we expected this independence would provide advantages in off-policy learning scenarios.

Our results suggest that critic networks can be updated through two pathways: via actor estimations and directly from environment feedback. In environments with dense, immediate rewards (like CartPole, MountainCar, and Pendulum), the direct feedback pathway appears sufficient to drive learning even with random actions. The critic can gradually build meaningful representations by associating state-action pairs with observed rewards and subsequent states. However, this mechanism becomes less effective as the temporal distance between actions and their consequences increases, making it harder to propagate reward signals through the network.

Uniform random sampling also completely removes the exploration problem from the learning process. Rather than relying on an actor to discover high-value regions of the state-action space, random sampling guarantees coverage of the entire space. This explains why algorithms

converged faster and more consistently in off-policy training for simpler environments, and why even problematic cases for AFU in on-policy training (like Pendulum’s bottom position) were successfully learned in off-policy mode.

The LunarLander environment represents a significantly more complex control task requiring precise sequential actions to achieve successful landing. As shown in Fig. 4, none of the algorithms demonstrated meaningful learning progress under random sampling. This contrasts sharply with on-policy results (Fig. 18), where AFU and SAC successfully learned optimal policies, and even DDPG converged to a « compromise » policy that keeps the lander airborne to avoid crashes.




Fig. 4. – Evolution of training performance for SAC, AFU, and DDPG algorithms in the Lunar Lander environment under off-policy learning with uniformly sampled state-action pairs. The curves represent interquartile means of returns over 200,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. Performance evaluations were conducted every 100 training steps by deploying the current policy for 10 complete episodes.

It seems the direct update mechanism becomes ineffective in environments like LunarLander,

where rewards depend on long sequences of coordinated actions. In such environments, the probability of randomly sampling informative transitions becomes exceedingly small as the state and action dimensions increase. This illustrates the problem of reward rarity - using our maze analogy, if the rewarding state-action pairs represent only a tiny fraction of the overall space (e.g., 1 in  $10^6$  cases), the probability of discovering them through random sampling becomes negligible without prohibitively large amounts of data. Furthermore, even when such rare rewards are encountered, propagating their value through the network becomes challenging without specifically structured network architectures.

The failure of all algorithms in LunarLander indicates fundamental limitations to learning from purely random data when tasks require specific sequential decision-making. This limitation affects all tested approaches regardless of their specific mechanism for handling the maximum Q-value problem.

We conducted an additional experiment using pre-trained critic values from successful on-policy training to initialize the critics for off-policy learning in LunarLander. Even with this advantage, which eliminates the need to discover reward signals from scratch, none of the algorithms showed convergence. This suggests that the challenge isn’t merely in propagating values but in establishing useful connections between randomly sampled state-action pairs and the sequential decision-making required for successful landing. This result further highlights the incompleteness of our understanding of the max Q update mechanisms.

We also experimented with larger neural networks to test if increased capacity would help maintain and propagate rare reward signals, but these attempts did not yield successful convergence. This suggests that the challenge may not be solely about network capacity but about the fundamental information flow in off-policy learning with random sampling.



Our findings suggest that the theoretical advantage of AFU’s critic update independence does not necessarily translate to superior performance in off-policy scenarios with random sampling. For simple environments with dense reward signals, traditional actor-critic methods can learn effectively from random data despite their theoretical dependence on actor quality. For complex environments with sparse rewards dependent on action sequences, all current approaches struggle regardless of their update mechanisms.

These results indicate that the interdependence between actor and critic in traditional algorithms may be less restrictive for off-policy learning than theoretically predicted. The ability to learn from random samples appears more closely tied to the reward structure, state-action space dimensionality, and temporal dependency of the environment than to the specific algorithm design.

*Offline to online transition.* Our second experimental investigation focused on the transition from offline to online reinforcement learning. This scenario is particularly relevant in robotics applications, where pretraining an agent on a dataset before deployment can significantly improve sample efficiency and reduce costly or potentially dangerous interactions with the real environment.

Offline reinforcement learning involves training an agent using a fixed dataset of transitions without any direct environment interaction. This dataset often contains samples from expert or near-expert policies, enabling the agent to develop foundational value estimations and behaviors. When properly trained, the agent can then transition to online learning, where it interacts directly with the environment to refine and optimize its policy further.

A significant challenge in this transition is the distribution shift problem. When an agent trained on a fixed dataset begins interacting with the environment, it encounters states and actions not represented in the original dataset. Traditional approaches address this problem by adopting

conservative policies that underestimate Q-values for out-of-distribution actions, preventing the agent from selecting potentially harmful actions.

However, this conservatism can lead to performance degradation during the transition to online learning. For algorithms that severely underestimate Q-values, actions encountered during online exploration may appear disproportionately valuable compared to those in the dataset. This can trigger a cycle where the agent repeatedly selects suboptimal actions it perceives as valuable, leading to significant performance drops. This problem is particularly severe when the offline dataset comes from an expert policy, as most out-of-distribution actions are likely suboptimal.

Unlike conservative methods, AFU addresses the distribution shift problem through its advantage function decomposition approach. By not artificially constraining Q-values for unseen actions, but instead directly estimating maximum Q-values through its value-advantage decomposition, AFU theoretically should maintain more stable performance during the transition from offline to online learning, avoiding the sharp performance degradation often observed with other algorithms.

To test this hypothesis, we created datasets for both Pendulum and Lunar Lander environments. These datasets were collected during standard training runs, capturing transitions during performance evaluation episodes. This approach provided a mixture of near-optimal trajectories and more diverse actions, creating a balanced dataset for offline learning. Using these datasets, algorithms were trained for 200,000 time steps in offline mode before transitioning to online learning for an additional 200,000 steps.

The Pendulum environment yielded unexpected results (Fig. 5). All algorithms achieved relatively good performance during offline learning, with SAC nearly converging to optimal policy. This surprising success of SAC in the offline phase

aligns with our earlier observation that SAC can learn effectively from off-policy data in this environment.




Fig. 5. – Evolution of training performance for AFU, SAC, IQL, and CALQL algorithms in the Pendulum environment during offline-to-online transition learning. The curves represent interquartile means of returns over 400,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. The vertical dashed line at 200,000 steps indicates the transition point from offline learning (using a pre-collected dataset) to online learning (with direct environment interaction).

IQL maintained stable performance during the transition, without the expected decline. This stability may be attributed to the simplicity of the Pendulum environment, where the network learned during offline training generalized well to the online phase. CALQL, being a conservative algorithm, showed limited progress during offline learning, consistent with its pessimistic approach. AFU demonstrated incomplete convergence during offline learning, which corresponds to its difficulties in on-policy learning for this environment, particularly in estimating values when the pendulum is in the bottom position.

The visualization of value functions (Figures Fig. 28, Fig. 29, Fig. 30, Fig. 31) confirms these observations, revealing similar value estimation patterns to those observed in on-policy learning.

The Lunar Lander environment provided more insightful results (Fig. 6). SAC nearly converged to a compromise « hovering » policy during offline learning, then rapidly improved toward the optimal landing policy after transitioning to online learning.

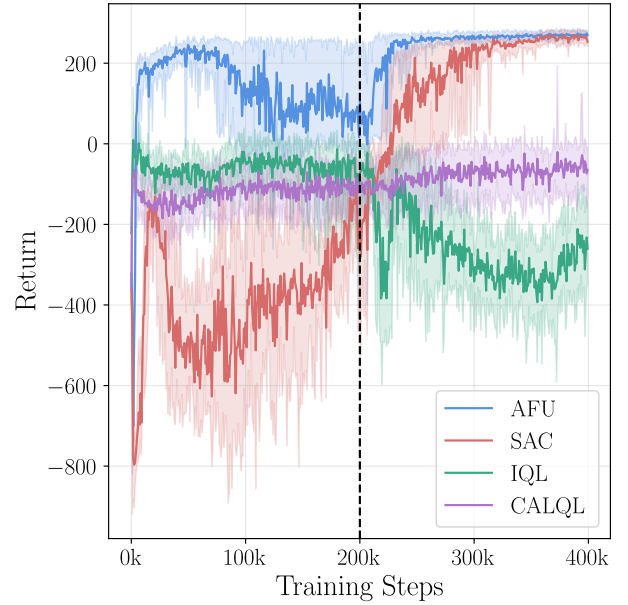


Fig. 6. – Evolution of training performance for AFU, SAC, IQL, and CALQL algorithms in the Lunar Lander environment during offline-to-online transition learning. The curves represent interquartile means of returns over 400,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. The vertical dashed line at 200,000 steps indicates the transition point from offline learning (using a pre-collected dataset) to online learning (with direct environment interaction).

IQL, as a pessimistic algorithm, never progressed beyond the compromise policy throughout both phases. Its conservatism prevented exploration of potentially risky but necessary landing strategies, keeping it within the boundaries of the safer hovering policy. As expected from a conservative

approach, its performance declined when transitioning to online learning.

CALQL similarly remained at the compromise policy level. Its calibration mechanism prevented the severe underestimation that leads to performance collapse, but its conservative nature still constrained exploration, resulting in negative interquartile mean rewards even after online learning, indicating unsuccessful landing attempts on average.

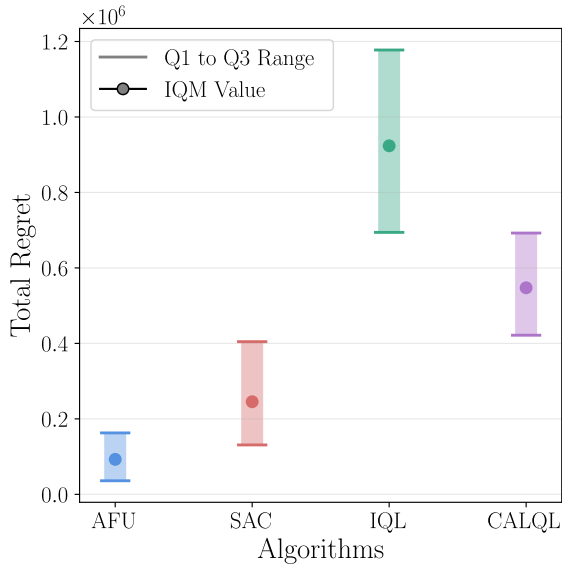


Fig. 7. – Comparison of total regret for different algorithms during offline-to-online transition in the Lunar Lander environment. Lower values indicate better performance, with vertical bars showing the Q1 to Q3 range and dark markers representing the interquartile mean (IQM) value across 5 independent trials.

AFU exhibited particularly interesting behavior for our hypothesis. It achieved high performance early in offline learning, followed by a slight decline, while maintaining predominantly successful landing policies as indicated by positive returns. When transitioning to online learning, AFU immediately resumed improvement, rapidly converging to the optimal policy. The interquartile range suggests consistent successful landings across trials, confirming effective learning. This pattern aligns with the ideal scenario for

robotics applications, where offline learning establishes fundamental competencies that are then refined through online interaction.

To quantitatively assess performance during transition, we measured the area under the curve representing the difference between each algorithm’s performance and the optimal value (Fig. 7). This total regret metric shows that AFU significantly outperformed other algorithms, with the smallest gap between its performance and the optimal policy. This confirms our hypothesis that AFU’s approach to solving the max-Q problem provides advantages during the offline-to-online transition.

## CONCLUSION

Our investigation into AFU’s capabilities across different learning scenarios has yielded several important insights regarding our initial hypotheses.

Regarding our first hypothesis about AFU’s ability to learn from randomly sampled data, the results were more nuanced than anticipated. In simpler environments (CartPole, MountainCar, and Pendulum), all algorithms, including traditional methods like SAC and DDPG, successfully learned from random data despite theoretical limitations in their critic update mechanisms. This success can be attributed to direct environment feedback enabling critic updates without relying on actor-proposed actions, particularly in environments with dense, immediate rewards. However, in the more complex Lunar Lander environment, which requires coordinated action sequences, no algorithm achieved meaningful learning from random sampling, including AFU. This highlights fundamental limitations in learning from purely random data when tasks require sequential decision-making, regardless of the specific max-Q estimation approach.

Our second hypothesis regarding AFU’s stability during offline-to-online transitions was supported by our experiments, particularly in the Lunar Lander environment. AFU demonstrated super-

ior performance during this transition, quickly adapting to online learning without the performance degradation observed in more conservative approaches like IQL and CALQL. The quantitative analysis of total regret confirmed AFU’s advantage, showing significantly smaller deviations from optimal performance compared to other algorithms. This validates our theoretical understanding that AFU’s approach to solving the max-Q problem, which does not artificially constrain Q-values for unseen actions, provides tangible benefits in realistic learning scenarios.

Overall, our work demonstrates that while the independence between critic and actor updates in AFU does not necessarily translate to superior performance in all off-policy learning scenarios, it provides significant advantages during the critical transition from offline to online learning. This finding has practical implications for applications where sample efficiency is paramount, such as robotics, where pretraining on datasets before deployment can reduce costly or potentially dangerous interactions with the environment. The results suggest that AFU’s unique architecture makes it particularly well-suited for such applications.

## APPENDIX A

### *Experimental Details.*

The custom environment wrappers used in this study include the following methods:

- `_set_state`: Allows resetting the environment to a specific state.
- `get_obs`: Retrieves the current observation.
- `get_observation_space`: Returns the observation space of the environment.
- `get_action_space`: Returns the action space of the environment.

These wrappers were implemented to ensure compatibility with the experiments.

### *Hyperparameters.*

The following hyperparameters were used for the experiments:

- Hidden size : 256
- Learning rate:  $3 \times 10^{-4}$
- Discount factor ( $\gamma$ ): 0.99
- Batch size: 256
- Replay buffer size: 200000
- Target network update rate ( $\tau$ ): 0.01
- Gradient reduction (for AFU only) : 0.5
- Noise standard deviation (for DDPG only) : 0.1
- Expectile regression ( $\tau'$ ) (for IQL only): 0.9
- Inverse temperature ( $\beta$ ) (for IQL only): 3.0
- Conservative factor ( $\alpha$ ) (for Cal-QL only): 5
- Random repetitions ( $n$ ) (for Cal-QL only): 4

### *Reproducibility.*

The code for all algorithms and experiments is available at <https://github.com/paulchambaz/afu-study>. The repository includes detailed instructions for setting up the environment and running the experiments. The code is licensed under the GPLv3 license and freely available.

## APPENDIX B

### *Implementation Notes on CAL-QL.*

In our review of the CALQL implementation presented in the original paper, we identified two minor inconsistencies in the provided pseudocode that we corrected in our implementation:

1. In the target Q-value calculation, the original code shows:

```
target_qval =
target_critic(batch['observations'],
next_pi_actions)
```

However, this should use the next observations according to the Bellman equation:

```
target_qval =
target_critic(batch['next_observations'],
next_pi_actions)
```

2. In the random actions Q-value calculation, the original code indicates:

```
q_rand_is = critic(batch['observations'],
random_actions) - random_pi
```

For consistency with the log-probabilities used elsewhere in the algorithm, this should instead use the logarithm of the uniform density:

```
q_rand_is = critic(batch['observations'],
random_actions) - log(random_pi)
```

These corrections align with both the mathematical principles of the algorithm and the authors' actual implementation in their source code repository. While these inconsistencies are minor and likely typographical in nature, they are worth noting for those aiming to implement CALQL correctly.

#### *Implementation Notes on IQL.*

In our review of the Implicit Q-Learning implementation presented in the original paper, we identified a minor inconsistency in the provided formula that we corrected in our implementation:

The policy loss function in Equation (7) of the paper shows:

$$L_{\pi}(\phi) = \mathbb{E}_{(s,a) \sim D} [\exp(\beta(Q_{\theta}(s,a) - V_{\psi}(s))) \log \pi_{\phi}(a | s)]$$

However, since this is a loss function that should be minimized during optimization, the correct formula should include a negative sign:

$$L_{\pi}(\phi) = \mathbb{E}_{(s,a) \sim D} [-\exp(\beta(Q_{\theta}(s,a) - V_{\psi}(s))) \log \pi_{\phi}(a | s)]$$

This correction aligns with the mathematical principles of advantage-weighted regression, where the objective is to maximize the likelihood of actions with high advantages. While this inconsistency is minor and likely typographical in nature, it's worth noting for those aiming to implement IQL correctly.

## APPENDIX C

#### *Implementation detail of DDPG.*

DDPG optimizes a continuous control policy using deterministic actions and off-policy learning. It uses two neural networks with parameters  $\theta$  and  $\varphi$ : a critic (Q-network) and an actor

(policy network). The critic  $Q_{\theta}(s,a)$  estimates expected returns for state-action pairs, while the actor  $\mu_{\varphi}(s)$  maps states directly to deterministic actions. Target networks with parameters  $\theta'$  and  $\varphi'$  stabilize learning through soft updates with rate  $\tau$ . Exploration during training uses Gaussian noise  $\mathcal{N}(0, \sigma^2)$  with standard deviation  $\sigma$ . Experiences  $(s, a, r, s')$  are stored in a replay buffer  $\mathcal{B}$  of size  $N$ , from which we sample mini-batches of size  $B$  for updates.

The first is a Q-network, which does a forward pass on a state-action pair and estimates the expected reward that the agent gets after doing a given action in a given state. To update its weights, we compute the following loss:

$$L_Q(\theta) = \mathbb{E}_{(s,a,r,s') \sim B} \left[ \left( Q_{\theta}(s,a) - \left( r + \gamma Q_{\theta'}(s', \mu_{\varphi'}(s')) \right) \right)^2 \right]$$

The second is a policy network (Actor), which outputs deterministic actions for any given state. Unlike SAC's stochastic policy, it directly maps states to optimal actions without probability distributions. Using a tanh activation to bound actions within  $[-1, 1]$ , the policy is optimized by maximizing the expected Q-value:

$$L_{\mu}(\varphi) = -\mathbb{E}_{s \sim B} [Q_{\theta}(s, \mu_{\varphi}(s))]$$

#### *Implementation detail of SAC.*

SAC implements an actor-critic method with entropy regularization for improved exploration. It uses five neural networks: twin critics with parameters  $\theta_1$  and  $\theta_2$  to reduce overestimation bias, a value network with parameters  $\psi$ , a target value network with parameters  $\psi'$ , and a policy network with parameters  $\varphi$ . The temperature parameter  $\alpha$  controls exploration, with target entropy set to  $-|\mathcal{A}|$  where  $\mathcal{A}$  is the action space. Experience tuples  $(s, a, r, s')$  are stored in a replay buffer  $\mathcal{B}$  of capacity  $N$ , from which mini-batches of size  $B$  are sampled. Networks are updated using learning rates  $l_Q$ ,  $l_V$ ,  $l_{\pi}$ , and  $l_{\alpha}$  for critics, value network, policy, and temperature respectively, with target networks soft-updated at rate  $\tau$ .

The first is a V-network, which does a forward pass on a given state and estimates the expected reward an agent can gain in the state. To update its weights, we compute the following loss:

$$L_V(\psi) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{B}} \left[ \left( V_\psi(s) - \left( \min Q_{\theta_i}(s, a_s) - \alpha \log(\pi_\varphi(a_s | s)) \right) \right)^2 \right]$$

where  $a_s \sim \pi_\varphi(\cdot | s)$  is an action sampled from the policy.

The second is a Q-network, which does a forward pass on a state-action pair and estimates the expected reward that the agent gets after doing a given action in a given state. To update its weights, we compute the following loss:

$$L_Q(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{B}} \left[ \left( Q_{\theta_i}(s, a) - r - \gamma V_{\psi^{\text{target}}}(s') \right)^2 \right]$$

The last is a policy network, which parameterizes a Gaussian distribution over actions. It outputs mean actions and learns log standard deviations as parameters. Using the reparameterization trick and tanh squashing for bounded actions, the policy is optimized by minimizing:

$$L_\pi(\phi) = \mathbb{E} \left[ \alpha \log(\pi_\phi(a_s | s)) - Q_{\theta_i}(s, a_s) \right]$$

*Implementation detail of AFU.* Actor Free critic Updates implements an actor-critic architecture with a novel advantage decomposition approach. It uses six neural networks: a main critic network with parameters  $\psi$ , twin value networks with parameters  $\varphi_1$  and  $\varphi_2$ , their target versions with parameters  $\varphi_{1'}$  and  $\varphi_{2'}$ , twin advantage networks with parameters  $\xi_1$  and  $\xi_2$ , and a policy network with parameters  $\theta$ . The temperature parameter  $\alpha$  controls exploration, with target entropy set to  $-|\mathcal{A}|$  where  $\mathcal{A}$  is the action space. The gradient reduction parameter  $\rho \in [0, 1]$  controls value function learning dynamics. Experiences  $(s, a, r, s')$  are stored in a replay buffer  $\mathcal{B}$  of capacity  $N$ , from which mini-batches of size  $B$  are sampled. Networks are updated using learning rates  $l_Q$ ,  $l_V$ ,  $l_{\psi}$ , and  $l_\alpha$  for critic, value/advantage networks, policy, and temperature

respectively, with target networks soft-updated at rate  $\tau$ .

The first is a Q-network, which does a forward pass on a state-action pair and estimates the expected reward of an agent taking action  $a$  in state  $s$ . To update its weights, we compute the following loss:

$$L_Q(\psi) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{B}} \left[ \left( Q_\psi(s, a) - r - \gamma \min_{i \in \{1,2\}} V_{\varphi_{i'}}(s') \right)^2 \right]$$

The second is a combined value and advantage loss for each network pair, which updates both value and advantage networks based on the relation  $Q(s, a) = V(s) + A(s, a)$ . To update their weights, we compute:

$$L_{VA}(\varphi_i, \xi_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{B}} \left[ Z \left( \Upsilon_i^a(s) - r - \gamma \min_{i \in \{1,2\}} V_{\varphi_{i'}}(s'), A_{\xi_i}(s, a) \right) \right]$$

where:

$$Z(x, y) = \begin{cases} (x + y)^2 & \text{if } x \leq 0 \\ x^2 + y^2 & \text{otherwise} \end{cases}$$

$$\Upsilon_i^a(s) = (1 - \rho \cdot I_i(s, a)) V_{\varphi_i}(s) + \rho \cdot I_i(s, a) \cdot V_{\varphi_i}^{\text{nograd}}(s)$$

$$I_i(s, a) = \begin{cases} 1 & \text{if } V_{\varphi_i}(s) + A_{\xi_i}(s, a) < Q_\psi(s, a) \\ 0 & \text{otherwise} \end{cases}$$

The last is a policy network that parameterizes a Gaussian distribution over actions. It outputs mean actions and learns log standard deviations as parameters. Using the reparameterization trick and tanh squashing for bounded actions, the policy is optimized by minimizing:

$$L_\pi(\theta) = \mathbb{E}_{s \sim \mathcal{B}} \left[ \alpha \log(\pi_\theta(a_s | s)) - Q_\psi(s, a_s) \right]$$

*Implementation detail of IQL.* Implicit Q-Learning is an offline reinforcement learning algorithm that decouples value learning from policy improvement. It uses five neural networks: twin critics with parameters  $\theta_1$  and  $\theta_2$ , their target versions with parameters  $\theta_{1'}$  and  $\theta_{2'}$ , a value network with parameters  $\psi$ , and a policy network

with parameters  $\varphi$ . The expectile regression parameter  $\tau' \in [0, 1]$  controls the conservatism in value estimation, while the temperature parameter  $\beta > 0$  determines how aggressively to exploit advantages. Experience tuples  $(s, a, r, s')$  are stored in a replay buffer  $\mathcal{B}$  of capacity  $N$ , from which mini-batches of size  $B$  are sampled. Networks are updated using learning rates  $l_Q$ ,  $l_V$ ,  $l_\pi$ , and  $l_\alpha$  for critics, value network, policy, and temperature respectively, with target networks soft-updated at rate  $\tau$ .

The first is a value network, which estimates the expected return from a state without considering actions. To update its weights, we compute the asymmetric  $L_2$  loss:

$$L_V(\psi) = \mathbb{E}_{(s,a) \sim \mathcal{B}} [L_2^{\tau'}(Q_\theta(s, a) - V_\psi(s))]$$

where  $L_2^{\tau'}(u) = |\tau' - \mathbb{1}(u < 0)| \cdot u^2$  is the asymmetric  $L_2$  loss with parameter  $\tau'$ .

The second is a Q-network, which estimates expected returns for state-action pairs. To update its weights, we compute the standard TD loss:

$$L_Q(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{B}} [(Q_{\theta_i}(s, a) - r - \gamma V_\psi(s'))^2]$$

The last is a policy network that parameterizes a Gaussian distribution over actions. It outputs mean actions and learns log standard deviations as parameters. Using the reparameterization trick and tanh squashing for bounded actions, the policy is optimized by maximizing likelihood weighted by advantages:

$$L_\pi(\varphi) = \mathbb{E}_{(s,a) \sim \mathcal{B}} [\exp(\beta \cdot \min(0, Q_{\theta'}(s, a) - V_\psi(s))) \cdot \log \pi_\varphi(a|s)]$$

where  $\beta$  controls the temperature and we clip advantages to be non-positive to avoid overoptimism.

*Implementation detail of CAL-QL.* Calibrated Q-learning extends offline RL with a conservative regularization approach. It uses three neural networks: a critic (Q-network) with parameters  $\psi$ , its target version with parameters  $\psi'$ , and a policy network with parameters  $\varphi$ . The tempe-

rature parameter  $\alpha > 0$  controls the conservative regularization strength, while the sampling parameter  $n$  determines how many random and policy-sampled actions are used for regularization. The algorithm incorporates Monte Carlo returns from a dataset as reference values  $V_\mu(s)$  to mitigate overestimation. Experience tuples  $(s, a, r, s')$  are stored in a replay buffer  $\mathcal{B}$  of capacity  $N$ , from which mini-batches of size  $B$  are sampled. Networks are updated using learning rates  $l_Q$ ,  $l_{\psi'}$ , and  $l_\alpha$  for critic, policy, and temperature respectively, with target networks soft-updated at rate  $\tau$ .

The first is a Q-network loss, which combines a standard TD loss with a conservative regularization term:

$$L_Q(\psi) = L_Q^{\text{TD}}(\psi) + \alpha \cdot L_Q^{\text{CON}}(\psi)$$

The TD loss follows the standard form:

$$L_Q^{\text{TD}}(\psi) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{B}} [(Q_\psi(s, a) - r - \gamma Q_{\psi'}(s', a_{s'}))^2]$$

The conservative regularization term penalizes Q-value overestimation:

$$L_Q^{\text{CON}}(\psi) = \mathbb{E}_{s \sim \mathcal{B}} \left[ \log \left( \sum_{i=1}^n \exp(Q_\psi(s, a_R^i) - \log(0.5^{|\mathcal{A}|})) \right) + \sum_{i=1}^n \exp(\max(V_\mu(s), Q_\psi(s, a_s^i) - \log \pi_\varphi(a_s^i|s))) - Q_\psi(s, a) \right]$$

where  $a_R^i$  are random actions sampled uniformly and  $V_\mu(s)$  are reference values from Monte Carlo returns.

The policy network is optimized through the standard SAC policy loss:

$$L_\pi(\varphi) = \mathbb{E}_{s \sim \mathcal{B}} [\alpha \log(\pi_\varphi(a_s | s)) - Q_\psi(s, a_s)]$$

## APPENDIX D

### *Off-policy learning experimental results.*




Fig. 8. – Distribution of returns for SAC, AFU, and DDPG algorithms in the CartPole environment under off-policy learning with uniformly sampled state-action pairs. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm.

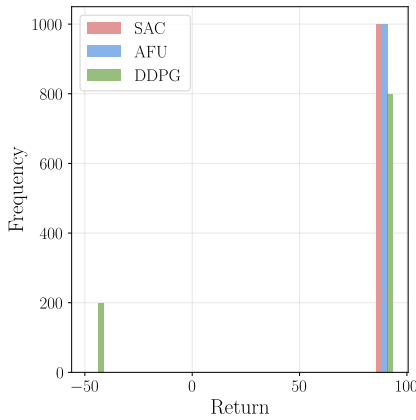


Fig. 9. – Distribution of returns for SAC, AFU, and DDPG algorithms in the Mountain Car environment under off-policy learning with uniformly sampled state-action pairs. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm.




Fig. 10. – Distribution of returns for SAC, AFU, and DDPG algorithms in the Pendulum environment under off-policy learning with uniformly sampled state-action pairs. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm.




Fig. 11. – Distribution of returns for SAC, AFU, and DDPG algorithms in the Lunar Lander environment under off-policy learning with uniformly sampled state-action pairs. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm.

### *On-policy learning experimental results.*



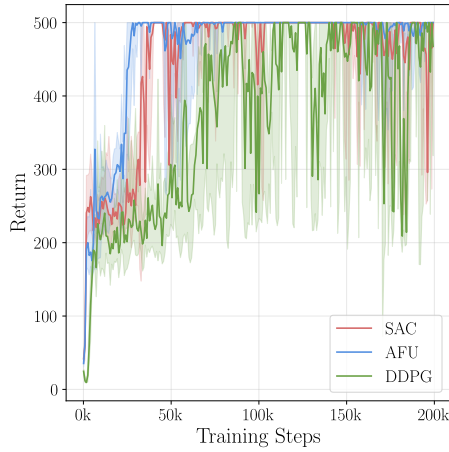


Fig. 12. – Evolution of training performance for SAC, AFU, and DDPG algorithms in the CartPole environment under on-policy learning with policy-generated samples. The curves represent interquartile means of returns over 200,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. Performance evaluations were conducted every 100 training steps by deploying the current policy for 10 complete episodes.




Fig. 13. – Distribution of returns for SAC, AFU, and DDPG algorithms in the CartPole environment under on-policy learning with policy-generated samples. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm.

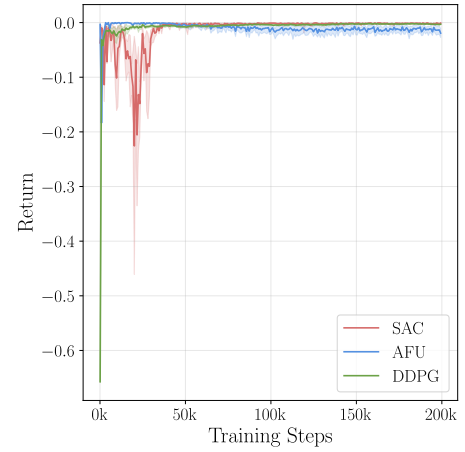


Fig. 14. – Evolution of training performance for SAC, AFU, and DDPG algorithms in the Mountain Car environment under on-policy learning with policy-generated samples. The curves represent interquartile means of returns over 200,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. Performance evaluations were conducted every 100 training steps by deploying the current policy for 10 complete episodes.




Fig. 15. – Distribution of returns for SAC, AFU, and DDPG algorithms in the Mountain Car environment under on-policy learning with policy-generated samples. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm.




Fig. 16. – Evolution of training performance for SAC, AFU, and DDPG algorithms in the Pendulum environment under on-policy learning with policy-generated samples. The curves represent interquartile means of returns over 200,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. Performance evaluations were conducted every 100 training steps by deploying the current policy for 10 complete episodes.




Fig. 17. – Distribution of returns for SAC, AFU, and DDPG algorithms in the Pendulum environment under on-policy learning with policy-generated samples. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm.




Fig. 18. – Evolution of training performance for SAC, AFU, and DDPG algorithms in the Lunar Lander environment under on-policy learning with policy-generated samples. The curves represent interquartile means of returns over 200,000 training steps, with shaded regions showing interquartile ranges across 5 independent trials. Performance evaluations were conducted every 100 training steps by deploying the current policy for 10 complete episodes.




Fig. 19. – Distribution of returns for SAC, AFU, and DDPG algorithms in the Lunar Lander environment under on-policy learning with policy-generated samples. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm.

*Offline online transition learning experimental results.*




Fig. 20. – Distribution of returns for AFU, SAC, IQL, and CALQL algorithms in the Pendulum environment after completing the offline-to-online transition learning process. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm after online fine-tuning.




Fig. 21. – Distribution of returns for AFU, SAC, IQL, and CALQL algorithms in the Lunar Lander environment after completing the offline-to-online transition learning process. The histogram shows the frequency of return values achieved during the final 1% of training episodes across 5 independent trials, representing the converged performance of each algorithm after online fine-tuning.

*Visualization of pendulum learned value function and action policy.*




Fig. 22. – Visualization of AFU algorithm's learned value function (left) and action policy (right) in the Pendulum environment under off-policy learning.




Fig. 23. – Visualization of SAC algorithm's learned value function (left) and action policy (right) in the Pendulum environment under off-policy learning.




Fig. 24. – Visualization of DDPG algorithm's learned value function (left) and action policy (right) in the Pendulum environment under off-policy learning.




Fig. 25. – Visualization of AFU algorithm's learned value function (left) and action policy (right) in the Pendulum environment under on-policy learning.




Fig. 28. – Visualization of AFU algorithm's learned value function (left) and action policy (right) in the Pendulum environment after offline-to-online transition.




Fig. 26. – Visualization of SAC algorithm's learned value function (left) and action policy (right) in the Pendulum environment under on-policy learning.




Fig. 29. – Visualization of SAC algorithm's learned value function (left) and action policy (right) in the Pendulum environment after offline-to-online transition.




Fig. 27. – Visualization of DDPG algorithm's learned value function (left) and action policy (right) in the Pendulum environment under on-policy learning.




Fig. 30. – Visualization of IQL algorithm's learned value function (left) and action policy (right) in the Pendulum environment after offline-to-online transition.




Fig. 31. – Visualization of CALQL algorithm’s learned value function (left) and action policy (right) in the Pendulum environment after off-line-to-online transition.

## BIBLIOGRAPHIE

- [1] Ilya Kostrikov, Ashvin Nair, et Sergey Levine. 2021. Offline Reinforcement Learning with Implicit Q-Learning. Consulté à l'adresse <https://arxiv.org/abs/2110.06169>
- [2] Mitsuhiro Nakamoto, Yuexiang Zhai, Anikait Singh, Max Sobol Mark, Yi Ma, Chelsea Finn, Aviral Kumar, et Sergey Levine. 2024. Cal-QL: Calibrated Offline RL Pre-Training for Efficient Online Fine-Tuning. Consulté à l'adresse <https://arxiv.org/abs/2303.05479>
- [3] Nicolas Perrin-Gilbert. 2024. AFU: Actor-Free critic Updates in off-policy RL for continuous control. Consulté à l'adresse <https://arxiv.org/abs/2404.16159>
- [4] Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, et Daan Wierstra. 2019. Continuous control with deep reinforcement learning. Consulté à l'adresse <https://arxiv.org/abs/1509.02971>
- [5] Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, et Sergey Levine. 2018. Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor. Consulté à l'adresse <https://arxiv.org/abs/1801.01290>